

Methods 3: Multilevel Statistical Modeling and Machine Learning

Week 39: *Link functions and fitting generalised linear multilevel models*
September 25, 2025

The course plan

Week 37:

Lesson 0: What is it all about?

Class 0: Setting up R and Python — and recollection of the general linear model (Lau)

Week 38:

Lesson 1: Multilevel linear regression

Class 1: Modelling subject level effects – and how do they differ from group level effects? (Lydia)

Week 39:

Lesson 2: Beyond linear regression — link functions

Class 2: What to do when the response variable is not continuous? (Lau)

Week 40:

Lesson 3: Evaluating generalised linear mixed models — and the spectre of collinearity

Class 3: Code review (Lydia)

Week 41:

Lesson 4: Explanation versus prediction

Class 4: Summarising mixed models (Lydia)

Week 42:

no teaching (go fetch potatoes instead)

Week 43:

Lesson 5: Mid-way evaluation and machine learning introduction

Class 5: Getting Python running (Lau)

Week 44:

Lesson 6: Linear and logistic regression revisited (machine learning)

Class 6: Categorizing responses based on informed guesses (Lau)

Week 45:

Lesson 7: Model evaluation and hyperparameter tuning

Class 7: Code review (Lydia: To be moved)

Week 46:

Lesson 8: Dimensionality reduction — principled component analysis (PCA)

Class 8: What to do with very rich data? (Lau)

Week 47:

Lesson 9: Neural networks

Class 9: Code review (Lau)

Week 48:

Lesson 10: Unsupervised classification

Class 10: Create video explainers of central concepts (shared resource) (Lydia)

Week 49:

Lesson 0 again: What was it all about? and final evaluation

Class 11: Summarising machine learning (Lydia)

Week 50:

Student presentations:

Class 12: Work on portfolios: Ask anything! (Lydia & Lau)

Recap – week 02

- Design matrices
 - reflect the choices we make in modelling
 - sometimes, our choices cannot be fit
 - the variance-covariance matrix will sometimes reveal why
- Partial pooling
 - allows for subject weighting or average weighting on a case-by-case basis, depending on the reliability of information available

Learning goals and outline

Link functions and fitting generalised linear multilevel models

- 1) Understanding that we can extend the scope of our general linear model by using appropriate link functions and data distributions
- 2) Understanding that we can extend this scope to multilevel modelling as well
- 3) Getting an understanding of maximum likelihood estimation

Extending the scope of the general linear model by using appropriate link functions and data distributions

With a focus on logistic regression (*Booleans*) and Poisson regression (*counts*)

Generalised linear model

At least four ingredients needed

- 1) A data vector: $y = (y_1, \dots, y_n)$
- 2) Predictors: X and coefficients β , forming a linear predictor $X\beta$
- 3) A *link function* g : yielding a vector of transformed data $\hat{y} = g^{-1}(X\beta)$
that are used to model the data
- 4) A data distribution: $p(y|\hat{y})$

$$(X\beta = \beta_0 + X_1\beta_1 + \dots + X_k\beta_k)$$

Gelman A, Hill J (2006) Data Analysis Using Regression and Multilevel/Hierarchical Models. Cambridge University Press: Chapter 6

Logistic regression



1) A data vector : $y = (y_1, \dots, y_n)$


```
data('mtcars')  
print(y <- mtcars$am)
```

```
## [1] 1 1 1 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 1 1 1 0  
0 0 0 0 1 1 1 1 1 1 1
```

```
print(n <- length(y))
```

```
## [1] 32
```


1) A data vector : $y = (y_1, \dots, y_n)$



2) Predictors: X and coefficients β , forming a linear predictor $X\beta$

```
X <- matrix(nrow=n, ncol=2) ## choosing two predictors
X[, 1] <- 1 ## modelling an intercept
X[, 2] <- mtcars$wt ## modelling a linear predictor based on weight
library(fields)
image.plot(x=1:dim(X)[2], y=1:dim(X)[1], z=t(X),
          xlab='Predictors', ylab='Observations', xaxt='n',
          main='Design matrix for logistic regression',
          legend.lab='Weight (lb/1000)')
axis(1, 1:dim(X)[2], c('Intercept', 'Slope'))
```

$X =$




$\beta =$ to be estimated


3) A link function g : yielding a vector of transformed data $\hat{y} = g^{-1}(X\beta)$ that are used to model the data

```
g      <- function(x) log(x / (1 - x)) ## logit
inv.g  <- function(x) exp(x) / (1 + exp(x)) ## logit-1
```


```
y <- seq(0.0001, 0.9999, 0.0001)
x <- 1:length(y)
plot(x, y, main='Probability of an event happening', xlab='Event no.',
     ylab='Probability of event happening')
```



```
plot(y, g(y), xlab='Probability of event', ylab='Logit transformed',  
     main='Logit transformed data\nRange from -Inf to Inf and continuous')
```



```
plot(x, inv.g(g(y)), main='Back on the original scale',  
     xlab='Event no.', ylab='Probability of event happening')
```



4) A data distribution: $p(y|\hat{y})$

$$\text{PMF}_{\text{Bernoulli}} = p^k (1-p)^{1-k} = \begin{cases} p & \text{if } k = 1, \\ q = 1-p & \text{if } k = 0. \end{cases}$$

$$p \in [0, 1]; k \in \{0, 1\}$$

```
dbernoul <- function(p, k) p^k * (1 - p)^(1 - k)


x <- c(0, 1)
pmf <- dbernoul(p=0.3, x)
par(font.lab=2, font.axis=2)
plot(x - 0.1, pmf, xaxt='n', xlab='Outcome', ylab='Probability', ylim=c(0, 1),
      xlim=c(-1.1, 2.1), main='Bernoulli probability mass function')
axis(side=1, at=c(0, 1), labels=c(0, 1))

pmf <- dbernoul(p=0.55, x)
points(x, pmf, col='red')

pmf <- dbernoul(p=0.88, x)
points(x + 0.1, pmf, col='blue')

legend('topright', pch=1, col=c('black', 'red', 'blue'),
      legend=c('p=0.3', 'p=0.55', 'p=0.88'), text.font=2)
```

Bernoulli probability mass function



Bringing it all together

Call: `glm(formula = am ~ wt, family = "binomial", data = mtcars)`

Coefficients:

(Intercept)	wt	$\hat{\beta}$
12.040	-4.024	

Degrees of Freedom: 31 Total (i.e. Null); 30 Residual

Null Deviance: 43.23

Residual Deviance: 19.18 AIC: 23.18

```
log.reg      <- glm(am ~ wt, data=mtcars, family='binomial')
y            <- log.reg$y
X            <- model.matrix(log.reg)
beta.hat     <- coef(log.reg)
lin.pred     <- X %*% beta.hat
y.hat        <- inv.g(lin.pred) # identical to fitted(log.reg)
likelihoods  <- dbinom(y, size=1, prob=y.hat) # bernoulli
likelihood   <- prod(likelihoods)
log.likelihood <- log(likelihood) ## check for yourself with logLik(log.reg)
res.deviance <- -2 * log.likelihood # null deviance is logLik(glm(am ~ 1))
n.predictors <- length(beta.hat)
AIC          <- 2 * n.predictors + res.deviance
print(log.reg)
```

live coding

4) A data distribution: $p(y | \hat{y})$

$p(y | \hat{y})$: the likelihood of observing y
when our fitted value is $\hat{y}$


$$\hat{y} = p(y=1)$$

$$p(y | \hat{y}) = \hat{y}^y (1 - \hat{y})^{1-y} : \text{Bernoulli distribution}$$
$$y \in \{0, 1\}; \hat{y} \in (0, 1)$$

```

x.pred <- seq(0, max(mtcars$wt), 0.01)
y.pred <- inv.g(coef(log.reg)[1] + coef(log.reg)[2] * x.pred)
plot(am ~ wt, data=mtcars, ylab='Probability of having manual transmission',
      , xlab='Weight (lb/1000)',
      yaxt='n', main='Fitted relationship between weight and transmission',
      ylim=c(-0.0, 1.0))
lines(x.pred, y.pred, lwd=3)

```



Why we need a data distribution

WE NEED A LIKELIHOOD TO MAXIMISE

$$L(\hat{\beta}|y) = \prod_{i=1}^n \hat{y}_i^{y_i} (1 - \hat{y}_i)^{(1-y_i)}$$

$$l(\hat{\beta}|y) = \sum_{i=1}^n y_i \log(\hat{y}_i) + \log(1 - \hat{y}_i)(1 - y_i)$$

Why is the log-likelihood preferred?

Poisson regression




1) A data vector : $y = (y_1, \dots, y_n)$

```
data('mtcars')  
print(y <- mtcars$carb)
```

```
## [1] 4 4 1 1 2 1 4 2 2 4 4 3 3 3 4 4 4 1 2 1 1 2 2 4 2 1 2 2 4 6 8 2
```

```
print(n <- length(y))
```


Relationship between weight and number of carburetors



2) Predictors: X and coefficients β , forming a linear predictor $X\beta$

```
X <- matrix(nrow=n, ncol=2) ## choosing two predictors
X[, 1] <- 1 ## modelling an intercept
X[, 2] <- mtcars$wt ## modelling a linear predictor based on weight
library(fields)
image.plot(x=1:dim(X)[2], y=1:dim(X)[1], z=t(X),
          xlab='Predictors', ylab='Observations', xaxt='n',
          main='Design matrix for Poisson regression',
          legend.lab='Weight (lb/1000)')
axis(1, 1:dim(X)[2], c('Intercept', 'Slope'))
```

$X =$




$\beta =$ to be estimated


3) A link function g : yielding a vector of transformed data $\hat{y} = g^{-1}(X\beta)$ that are used to model the data

```
g      <- function(x) log(x)
inv.g  <- function(x) exp(x)
```


```
y <- seq(0, 1000, 1)
x <- 1:length(y)
plot(x, y, main='Count of an event', xlab='Event no.',
     ylab='Event counts')
```



```
plot(y, g(y), xlab='Count of an event', ylab='Log transformed',  
     main='Log transformed data\nRange from -Inf to Inf and continuous')
```




```
plot(y, inv.g(g(y)), main='Back on the original scale',  
     xlab='Event no.', ylab='Count of an event')
```




4) A data distribution: $p(y|\hat{y})$

Poisson(λ)

$\lambda \in [0, \infty)$: the Expected value (rate)

$k \in N_0$; how many events happened

$$PMF_{Poisson} = \frac{\lambda^k e^{-\lambda}}{k!}$$



Bringing it all together

Call: `glm(formula = carb ~ wt, family = "poisson", data = mtcars)`

Coefficients:

(Intercept) wt

0.2391 0.2386

Degrees of Freedom: 31 Total (i.e. Null); 30 Residual


Null Deviance: 27.04

Residual Deviance: 21.96 AIC: 115.3

```
pois.reg <- glm(carb ~ wt, data=mtcars, family='poisson')
y <- pois.reg$y
X <- model.matrix(pois.reg)
beta.hat <- coef(pois.reg)
lin.pred <- X %*% beta.hat
y.hat <- inv.g(lin.pred) # identical to fitted(pois.reg)
likelihoods <- dpois(y, y.hat) # poisson
likelihood <- prod(likelihoods)
log.likelihood <- log(likelihood) ## check for yourself with logLik(pois.reg)
res.deviance <- 2 * sum(y * log(y/y.hat) - (y - y.hat))
n.predictors <- length(beta.hat)
AIC <- 2 * n.predictors - 2*log.likelihood
print(pois.reg)
```

live coding

```
x.pred <- seq(0, max(mtcars$wt), 0.01)
y.pred <- inv.g(coef(pois.reg)[1] + coef(pois.reg)[2] * x.pred)
plot(carb ~ wt, data=mtcars, ylab='Count of carburetors'
      , xlab='Weight (lb/1000)',
      main='Fitted relationship between weight and no. of carburetors',
      ylim=c(0.0, 10.0))
lines(x.pred, y.pred, lwd=3)
```



4) A data distribution: $p(y | \hat{y})$

$p(y | \hat{y})$: the likelihood of observing y
when our fitted value is $\hat{y}$

$$\hat{y} = \lambda$$

$$p(y | \hat{y}) = \frac{\hat{y}^y e^{-\hat{y}}}{y!} : \text{Poisson distribution}$$

$$y \in N_0; \hat{y} \in [0, \infty)$$

Why we need a data distribution

WE NEED A LIKELIHOOD TO MAXIMISE

$$L(\hat{\beta}|y) = \prod_{i=1}^n \frac{\hat{y}_i^{y_i} e^{-\hat{y}_i}}{y_i!}$$

$$l(\hat{\beta}|y) = \sum_{i=1}^n y_i \log(\hat{y}_i) - \hat{y}_i - \log(y_i!)$$

... the identity link

```
g      <- function(x) identity(x)
inv.g  <- function(x) identity(x) ##  $\text{identity}^{-1} = \text{identity}$ 
```


$$N(\mu, \sigma^2)$$

$$\mu \in \mathbb{R}$$

$$\sigma^2 \in \mathbb{R}_{>0}$$

$$x \in \mathbb{R}$$

$$\text{PDF} = (2\pi\sigma^2)^{-1/2} e^{-\frac{1}{2}\left(\frac{x-\mu}{\sigma}\right)^2}$$



Bringing it all together

Call: `glm(formula = mpg ~ wt, family = "gaussian", data = mtcars)`

Coefficients:

(Intercept) wt
 37.285 -5.344

Degrees of Freedom: 31 Total (i.e. Null); 30 Residual

Null Deviance: 1126

Residual Deviance: 278.3 AIC: 166

```
lin.reg      <- glm(mpg ~ wt, data=mtcars, family='gaussian')
y            <- lin.reg$y
n            <- length(y)
X            <- model.matrix(lin.reg)
beta.hat     <- coef(lin.reg)
lin.pred     <- X %*% beta.hat
y.hat        <- inv.g(lin.pred) # identical to fitted(lin.reg)
sigma.hat.squared <- 1/n * sum((y - y.hat)^2)
log.likelihood <- -n/2*log(2*pi) - n/2*log(sigma.hat.squared) - n/2
res.deviance  <- sum((y - y.hat)^2)
n.predictors  <- length(beta.hat) + 1 # sigma hat squared
AIC           <- 2 * n.predictors - 2*log.likelihood
print(lin.reg)
```


$$L(\beta, \sigma^2 | y) = \prod_{i=1}^n \frac{1}{\sqrt{2\pi\sigma^2}} e^{-\frac{(y_i - X\beta)^2}{2\sigma^2}}$$

$$l(\beta, \sigma^2 | y) = -\frac{n}{2} \log(2\pi) - \frac{n}{2} \log(\sigma^2) - \frac{1}{2\sigma^2} \sum_{i=1}^n (y_i - X\beta)^2$$

$$\hat{\sigma}^2 = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2$$

$$l(\hat{\beta}, \hat{\sigma}^2 | y) = -\frac{n}{2} \log(2\pi) - \frac{n}{2} \log(\hat{\sigma}^2) - \frac{n}{2}$$

Summary

- We can extend the General Linear Model to a Generalised Linear Model by ...
 - choosing appropriate link functions
 - we can define our own
 - and using appropriate data distributions
 - such that we have a likelihood function to maximise
- This can be used to extend the General Linear Multilevel Model to the Generalised Linear Multilevel Model
 - And we can get the benefits of partial pooling among other things

Maximising likelihood for multilevel models

Penalised least squares:

$$r^2(\theta, \beta, u) = \rho^2(\theta, \beta, u) + \|u\|^2 \quad (14)$$

$\rho^2(\theta, \beta, u)$ is the residual sum of squares

$$\mu_{Y|U=u} = X\beta + Z\Lambda_{\theta}u$$

(compare with: $\mu_Y = X\beta$ from classical regression)


($b = \Lambda_{\theta}u$) : u is estimating the variance between subjects

Example of how to fit


```
library(carData)
library(lme4)
inv.g <- function(x) exp(x) / (1 + exp(x)) ## logit-1
data('TitanicSurvival')
titanic <- TitanicSurvival

titanic$survived <- ifelse(titanic$survived == 'yes', 1, 0)

model.glm <- glm(survived ~ age, data=titanic, family='binomial')
```



```
model.glmm <- glmer(survived ~ age + (1 | passengerClass), data=titanic,  
                    family='binomial')
```



Poisson example

```
data('Wool')  
str(Wool)
```

```
## 'data.frame':   27 obs. of  4 variables:  
## $ len      : int  250 250 250 250 250 250 250 250 250 300 ...  
## $ amp      : int   8 8 8 9 9 9 10 10 10 8 ...  
## $ load     : int  40 45 50 40 45 50 40 45 50 40 ...  
## $ cycles: int  674 370 292 338 266 210 170 118 90 1414 ...
```

```
pois.glmm <- glmer(cycles ~ 1 + (1 | len) + (1 | amp) + (1 | load), data=Wool,  
                  family='poisson')
```

Learning goals and outline

Link functions and fitting generalised linear multilevel models

- 1) Understanding that we can extend the scope of our general linear model by using appropriate link functions and data distributions
- 2) Understanding that we can extend this scope to multilevel modelling as well
- 3) Getting an understanding of maximum likelihood estimation

The course plan

Week 37:

Lesson 0: What is it all about?

Class 0: Setting up R and Python — and recollection of the general linear model (Lau)

Week 38:

Lesson 1: Multilevel linear regression

Class 1: Modelling subject level effects – and how do they differ from group level effects? (Lydia)

Week 39:

Lesson 2: Beyond linear regression — link functions

Class 2: What to do when the response variable is not continuous? (Lau)

Week 40:

Lesson 3: Evaluating generalised linear mixed models — and the spectre of collinearity

Class 3: Code review (Lydia)

Week 41:

Lesson 4: Explanation versus prediction

Class 4: Summarising mixed models (Lydia)

Week 42:

no teaching (go fetch potatoes instead)

Week 43:

Lesson 5: Mid-way evaluation and machine learning introduction

Class 5: Getting Python running (Lau)

Week 44:

Lesson 6: Linear and logistic regression revisited (machine learning)

Class 6: Categorizing responses based on informed guesses (Lau)

Week 45:

Lesson 7: Model evaluation and hyperparameter tuning

Class 7: Code review (Lydia: To be moved)

Week 46:

Lesson 8: Dimensionality reduction — principled component analysis (PCA)

Class 8: What to do with very rich data? (Lau)

Week 47:

Lesson 9: Neural networks

Class 9: Code review (Lau)

Week 48:

Lesson 10: Unsupervised classification

Class 10: Create video explainers of central concepts (shared resource) (Lydia)

Week 49:

Lesson 0 again: What was it all about? and final evaluation

Class 11: Summarising machine learning (Lydia)

Week 50:

Student presentations:

Class 12: Work on portfolios: Ask anything! (Lydia & Lau)

Next time

- Overdispersion and other link functions
- Comparison of models
 - R^2
 - Log-likelihood testing
 - Akaike Information Criteria
 - Regularisation

Reading questions

- Chapter 14; focus on sections 14.1 and 14.2
 - Read for examples of advanced multilevel models and what they can be good for
- Chapter 15; focus on section 15.1
 - What is overdispersion?
- Suggestion: re-read pp. 31-42 of: Winter B (2013) Linear models and linear mixed effects models in R with linguistic applications. [arXiv:13085499](https://arxiv.org/abs/1308.5499) [cs]